

In this problem set, you will explore synchronization issues on a simulated multi-processor, shared-memory environment. We will not use threads-based programming, but instead will create an environment in which several UNIX processes share a memory region through mmap. Each process represents a parallel processor.

We will number each of these "virtual processors" with a small integer identifier which will be held in the global variable `my_procnum`. This is not the same as the UNIX process id, although you will probably need to keep track of the UNIX pids too. `my_procnum` ranges from 0 to `N_PROC-1`. `N_PROC` is the **maximum** number of virtual processors which your implementation is required to accept. For this project, define it as 64.

To implement sleeping and waking in this project, the UNIX signal facility will be used to simulate inter-processor interrupts. Use signal `SIGUSR1` and the system calls `sigsuspend` and `sigprocmask`, as discussed in class.

The starting point is an atomic test and set instruction. Since "some assembly is required," this will be provided to you in the file `tas.S` (32-bit), or `tas64.S` (64-bit). Use it with a makefile or directly with gcc, e.g. `gcc fifotest.c fifolib.c semlib.c tas.S` A .S file is a pure assembly language function. At the C level, it will work as:

```
int tas(volatile char *lock)
```

You will have to write your own prototype for the function since no header is provided. The `tas` function works as described in the lecture notes. A zero value means unlocked, and `tas` returns the *previous* value of `*lock`, meaning it returns 0 when the lock has been acquired, and 1 when it has not.

It is suggested that you implement a spin lock using this atomic TAS, and use that spin lock as a mutex to help implement the functions below. It will not be necessary to implement a full mutex lock with blocking/yielding, as that functionality will be built-in to the your semaphores.

### Problem 1 -- Implement semaphores

Create a module, called `sem.c`, with header file `sem.h`, which implements the semaphore operations defined below. You will need to make use of the spinlock mutex derived from the provided TAS function.

```
void sem_init(struct sem *s, int count);
```

Initialize the semaphore `*s` with the initial count. Initialize any underlying data structures. `sem_init` should only be called once in the program. If called after the semaphore has been used, results are unpredictable.

```
int sem_try(struct sem *s);
```

Attempt to perform the "P" operation (atomically decrement the semaphore). If this operation would block, return 0, otherwise return 1.

```
void sem_wait(struct sem *s);
```

Perform the P operation, blocking until successful. Blocking should be accomplished by noting within the `*s` that the current virtual processor needs to be woken up, and then sleeping using the `sigsuspend` system call until `SIGUSR1` is received. Assume that the extern int variable `my_procnum` exists and contains the virtual processor id of the caller. The implementation by which you keep track of waiting processors is up to you.

```
void sem_inc(struct sem *s);  
    Perform the V operation.  If any other processors were sleeping  
    on this semaphore, wake them by sending a SIGUSR1 to their  
    process id (which is not the same as the virtual processor number).
```

### Problem 2 -- A FIFO using semaphores

Now create a fifo module, `fifo.c` with associated header file `fifo.h`, which maintains a FIFO of unsigned longs using a shared memory data structure protected and coordinated **exclusively** with the semaphore module developed above. Depending on your approach you may or may not need to use all of the semaphore functions above. However, if your FIFO implementation takes more than about 100 lines of code, you are probably over-complicating things.

```
void fifo_init(struct fifo *f);  
    Initialize the shared memory FIFO *f including any  
    required underlying initializations.  fifo will  
    have a fifo length of MYFIFO_BUFSIZ elements, which should be  
    a static #define in fifo.h (a value of 4K is reasonable).  
void fifo_wr(struct fifo *f, unsigned long d);  
    Enqueue the data word d into the FIFO, blocking  
    unless and until the FIFO has room to accept it.  
unsigned long fifo_rd(struct fifo *f);  
    Dequeue the next data word from the FIFO and return it.  
    Block unless and until there are available words  
    queued in the FIFO.
```

### Problem 3 -- Test your FIFO

Create a framework for testing your FIFO implementation. Establish a `struct fifo` in shared memory and create two virtual processors, one of which will be the writer and the other the reader. Have the writer send a fixed number of sequentially-numbered data using `fifo_wr` and have the reader read these and verify that all were received.

Next, give your system the acid test by creating multiple writers, but one reader. In a successful test, all of the writers' streams will be received by the reader complete, in (relative) sequence, with no missing or duplicated items, and all processes will eventually run to completion and exit (no hanging). A suggested approach is to treat each datum (32-bit word) as a bitwise word consisting of an ID for the writer and the sequence number.

Use reasonable test parameters. Remember, an acid test of a FIFO where the buffer does not fill and empty quite a few times has a pH of 6.9, i.e. it isn't a very good acid. You should be able to demonstrate **failure** by deliberately breaking something in your implementation, e.g. reversing the order of two operations. You should then be able to demonstrate success under a variety of strenuous conditions.

*Submit all of the code comprising this final test system, i.e. your `sem.[ch]`, `fifo.[ch]` and `main.c` files, as well as output from your test program showing it ran correctly. If the output is very verbose, you may trim the uninteresting stuff with an appropriate annotation.*

*Your system should be bulletproof as far as locking and wait/wakeup in the face of multiple readers AND writers on the same FIFO (although you do not have to test the multiple readers). You will probably find that errors are quicker to appear on a true multi-processor system.*